LABORATORY RECORD

On

BIG DATA ANALYTICS

B.E (IT) – VII Sem

By

N.Durga Sai Lakshmi(160117737006)

Under the guidance of

Sri.K.Gangadhar Rao

Assistant Professor,

Dept. of IT,CBIT.

DEPARTMENT OF INFORMATION TECHNOLOGY

CHAITANYA BHARATHI INSTITUTE OF TECHNOLOGY (A)

(Affiliated to Osmania University; Accredited by NBA(AICTE) and NAAC(UGC), ISO Certified 9001:2015)

KOKAPET(V),GANDIPET(M), RR District HYDERABAD -75

Website: www.cbit.ac.in

2020-2021

160117737006
N.Durga Sai Lakshmi

Problem Statement-1: Write a Map-reduce application to find number of occurrences of each word from the given dataset.

Description:

In the MapReduce word count example, we find out the frequency of each word. Here, the role of Mapper is to map the keys to the existing values and the role of Reducer is to aggregate the keys of common values. So, everything is represented in the form of a Key-value pair.

Procedure:

$cd hadoop-3.2.1/

$cd sbin/

$cd start-all.sh

$jps

$cd..

$cd wc.py/

$hadoop fs -mkdir -p /wordcount/p1

$hadoop fs -copyFromLocal word.txt /wordcount/p1

$hadoop jar /home/hduser/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file wcmap.py -mapper wcmap.py -file wcred.py -reducer wcred.py -input /wordcount/p1 -output /wordcount/p1/output

$hadoop fs -cat /wordcount/p1/output/part-00000

Code:

wcmap.py

```
import sys

# input comes from STDIN (standard input)

for line in sys.stdin:

        # remove leading and trailing whitespace
```

```
        line = line.strip()

        # split the line into words

        words = line.split()

        # increase counters

        for word in words:

        # write the results to STDOUT (standard output);

        # what we output here will be the input for the

        # Reduce step, i.e. the input for reducer.py

        #

        # tab-delimited; the trivial word count is 1

                print '%s\t%s' % (word, 1)
```

wcred.py

```
from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None

# input comes from STDIN
for line in sys.stdin:
        # remove leading and trailing whitespace
        line = line.strip()

        # parse the input we got from mapper.py
        word, count = line.split('\t', 1)

        # convert count (currently a string) to int
        try:
```

160117737006
N.Durga Sai Lakshmi

```
        count = int(count)
        except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue

        # this IF-switch only works because Hadoop sorts map output
        # by key (here: word) before it is passed to the reducer
        if current_word == word:
        current_count += count
        else:
        if current_word:
        # write result to STDOUT
        print '%s\t%s' % (current_word, current_count)
        current_count = count
        current_word = word
if current_word == word:
        print '%s\t%s' % (current_word, current_count)
```

INPUT :

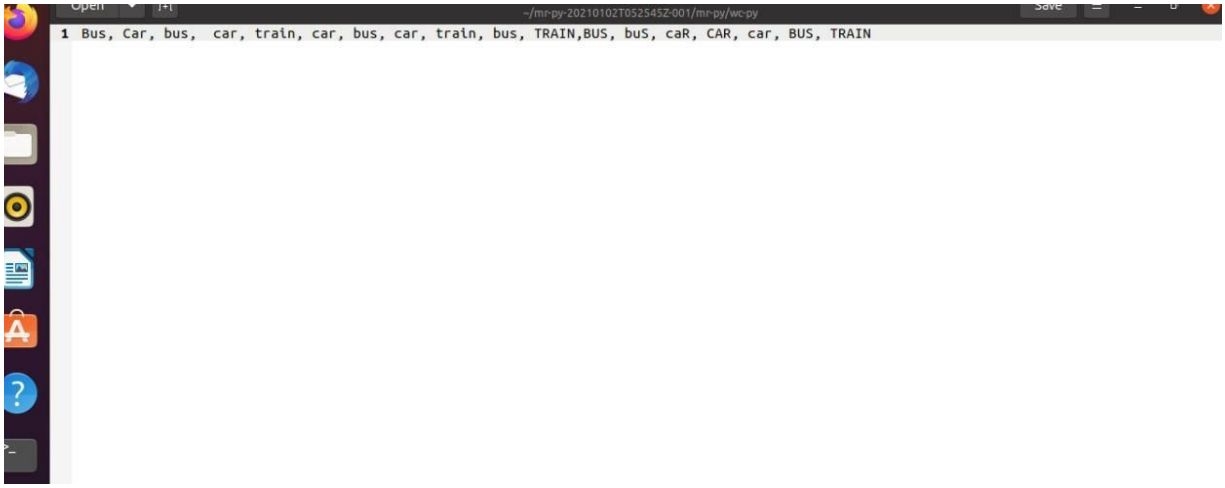

160117737006
N.Durga Sai Lakshmi

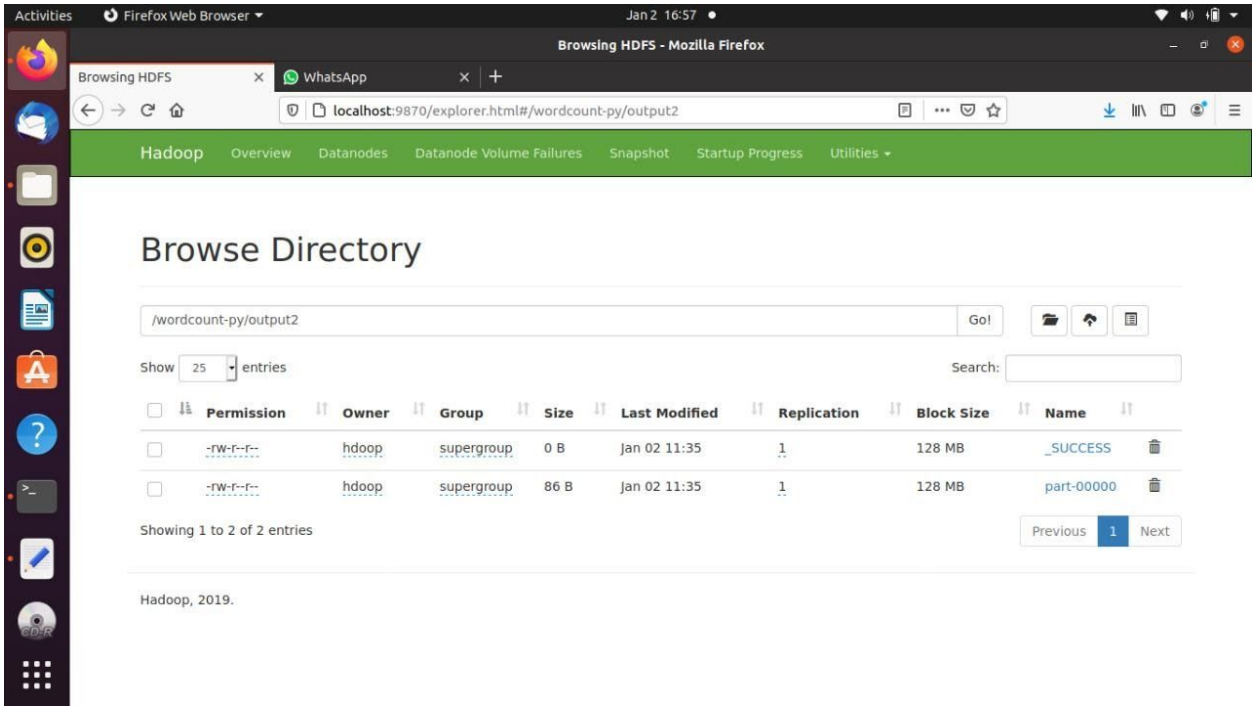

OUTPUT:

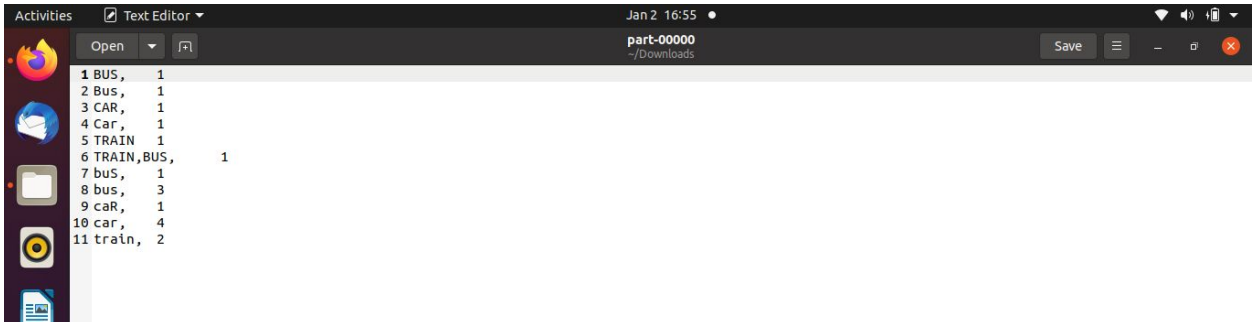

Problem Statement-2: Write a 3map-reduce application to predict maximum stock price for given dataset.

Description:

We are trying to find out the maximum closing price of each stock symbol. This means that we have to group the records by symbol so that we can calculate the maximum closing price by symbol. So we will output Stock Symbol as the key and close price as the value for each record. We now know what is going to be the Map's input and what is going to be the maps output.

Procedure:

$cd hadoop-3.2.1/

160117737006
N.Durga Sai Lakshmi

$cd sbin/

$cd start-all.sh

$jps

$cd..

$cd stock.py/

$hadoop fs -mkdir /stocksip

$hadoop fs -copyFromLocal stocks.txt /stocksip

$hadoop jar /home/hduser/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file stock-mapper.py -mapper stock-mapper.py -file stock-reducer.py -reducer stock-reducer.py -input /stocksip/stocks.txt -output /stocksout/output

$hadoop fs –cat stocksout/output/part-00000

Code:

Stock-mapper.py

```
import sys

for line in sys.stdin:
        line = line.strip()
        data = line.split(",")
        stock, price = data[1], data[6]
        print("%s\t%s"%(stock,price))
```

stock-reducer.py

```
import sys

max_price = 9999
max_stock = None

for line in sys.stdin:
        line = line.strip()
        stock,price = line.split("\t",1)
```

```
        if max_stock and max_stock!=stock:
                        if max_price > price:
                                max_price = price
                                max_stock = stock

        else:
                        max_stock, max_price = stock,max(max_price,price)

if max_stock:
        if max_price > price:
                                max_price = price
                                max_stock = stock

print("%s\t%s"%(max_stock,max_price))
```

OUTPUT:

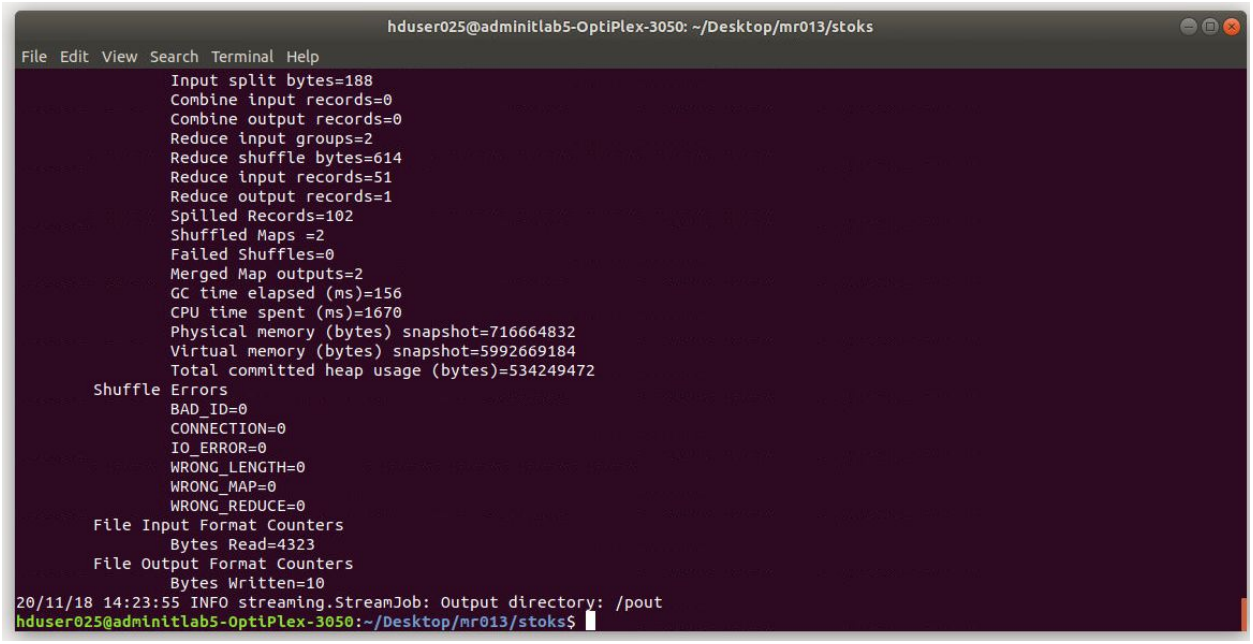

160117737006
N.Durga Sai Lakshmi

Problem Statement-3: Write a map-reduce application to find maximum temperature for a given year from the NCDC weather dataset.

Description:

MapReduce is based on set of key value pairs. So first we have to decide on the types for the key/value pairs for the input.

Map Phase: The input for Map phase is set of weather data files as shown in snap shot. The types of input key value pairs are *LongWritable* and Text and the types of output key value pairs are *Text* and *IntWritable*. Each Map task extracts the temperature data from the given year file. The output of the map phase is set of key value pairs. Set of keys are the years. Values are the temperature of each year.

Reduce Phase: Reduce phase takes all the values associated with a particular key. That is all the temperature values belong to a particular year is fed to a same reducer. Then each reducer finds the highest recorded temperature for each year. The types of output key value pairs in Map phase is same for the types of input key value pairs in reduce phase (*Text* and *IntWritable*). The types of output key value pairs in reduce phase is too *Text* and *IntWritable*.

Procedure:

$cd hadoop-3.2.1/

$cd sbin/

$cd start-all.sh

$jps

$cd..

$cd temp.py/

$ hadoop fs -mkdir /NCDCWeatherData

$hadoop fs -copyFromLocal NCDCWeatherData/* /NCDCWeatherData

$hadoop jar /home/hduser/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file temp-mapper.py -mapper temp-mapper.py -file temp-reducer.py -reducer temp-reducer.py -input /NCDCWeatherData -output /weatherout/output

$ hadoop fs –cat weatherout /output/part-00000

Code:
temp-mapper.py

```
import re
import sys
for line in sys.stdin:
val = line.strip()
(year, temp, q) = (val[15:19], val[87:92], val[92:93])
if (temp != "+9999" and re.match("[01459]", q)):
print "%st%s" % (year, temp)
```

temp-reducer.py

```
import sys
(last_key, max_val) = (None, 0)
for line in sys.stdin:
(key, val) = line.strip().split("t")
if last_key and last_key != key:
print "%st%s" % (last_key, max_val)
(last_key, max_val) = (key, int(val))
else:
(last_key, max_val) = (key, max(max_val, int(val)))
if last_key:
print "%st%s" % (last_key, max_val)
```

Expected Output:
1901 year has maximum temperature.

Problem Statement-4: Write a map-reduce application to find how many times a particular page has been accessed (use from the Apache Web Server log data).

Description:
In today's world the usage of internet has become very high and using all the logs from web server we can actually predict the customer moods in buying the product or can analyze the interests of Customer.

Procedure:

$cd hadoop-3.2.1/

160117737006
N.Durga Sai Lakshmi

$cd sbin/

$cd start-all.sh

$jps

$cd..

$cd web-log.py/

$hadoop fs -mkdir -p /weblogip

$hadoop fs -copyFromLocal test_access_log /weblogip
$hadoop jar /home/hduser/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file my_mapper_by_page.py -mapper my_mapper_by_page.py -file my_reducer.py -reducer my_reducer.py -input /weblogip/test_access_log -output /weblogout/output

$ hadoop fs –cat weblogout /output/part-00000

Code:
my_mapper_by_page.py

```
import sys

for line in sys.stdin:
        data = line.strip().split(" ")
        if len(data) == 10:
        page = data[6]
        print page
```

my_reducer.py

```
import sys

number = 0
oldKey = None

for line in sys.stdin:
        thisKey = line

   if oldKey and oldKey != thisKey:
        print oldKey, "\t", number
```

```
        oldKey = thisKey;
        number = 0

        oldKey = thisKey
        number += 1

if oldKey != None:
        print oldKey, "\t", number
```

OUTPUT:

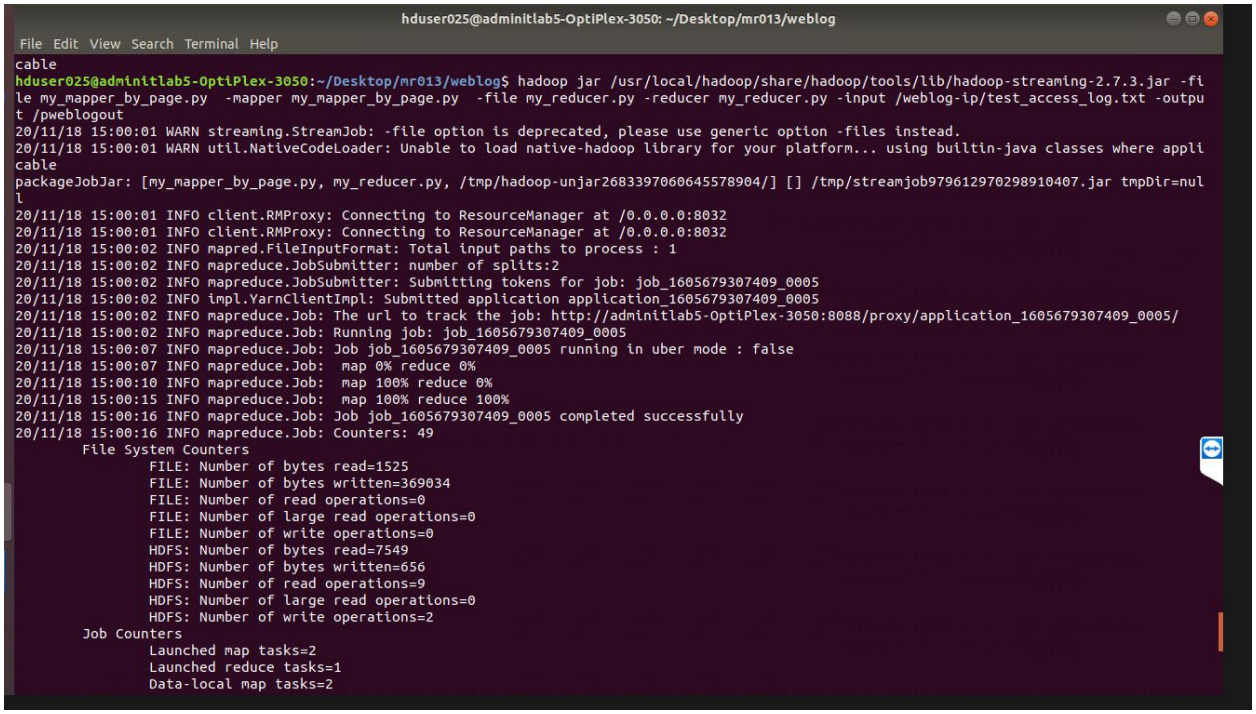

160117737006
N.Durga Sai Lakshmi

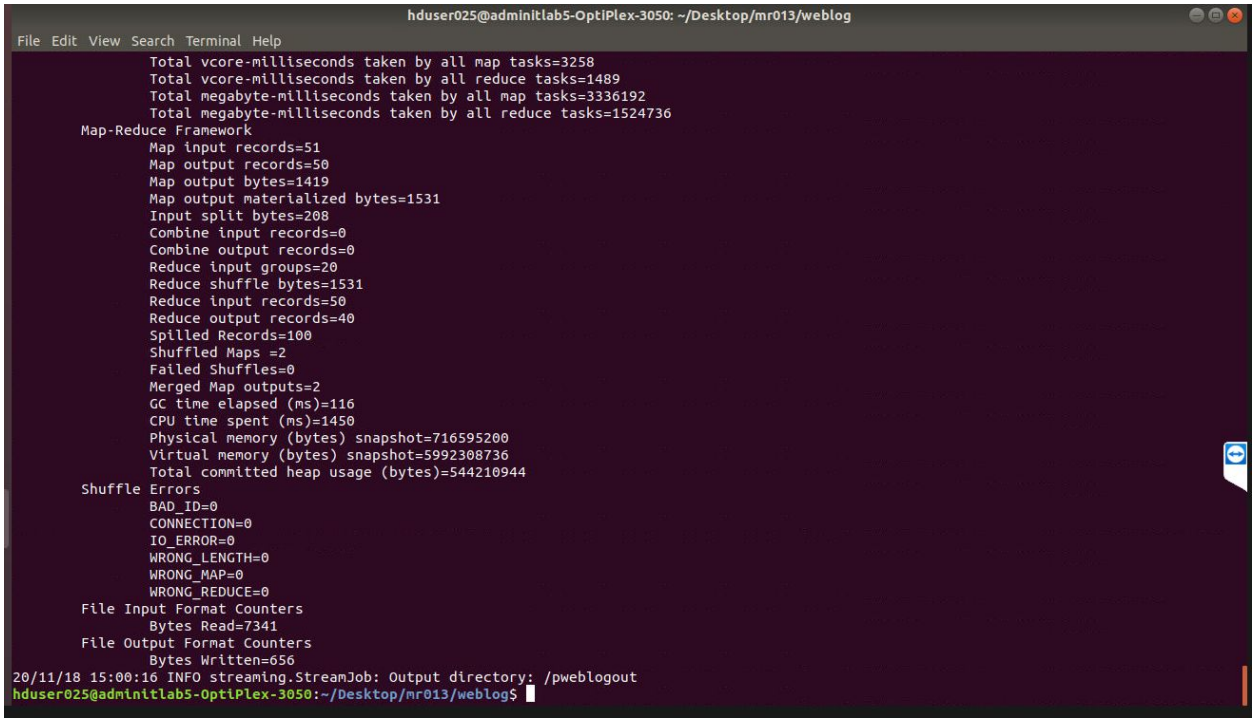

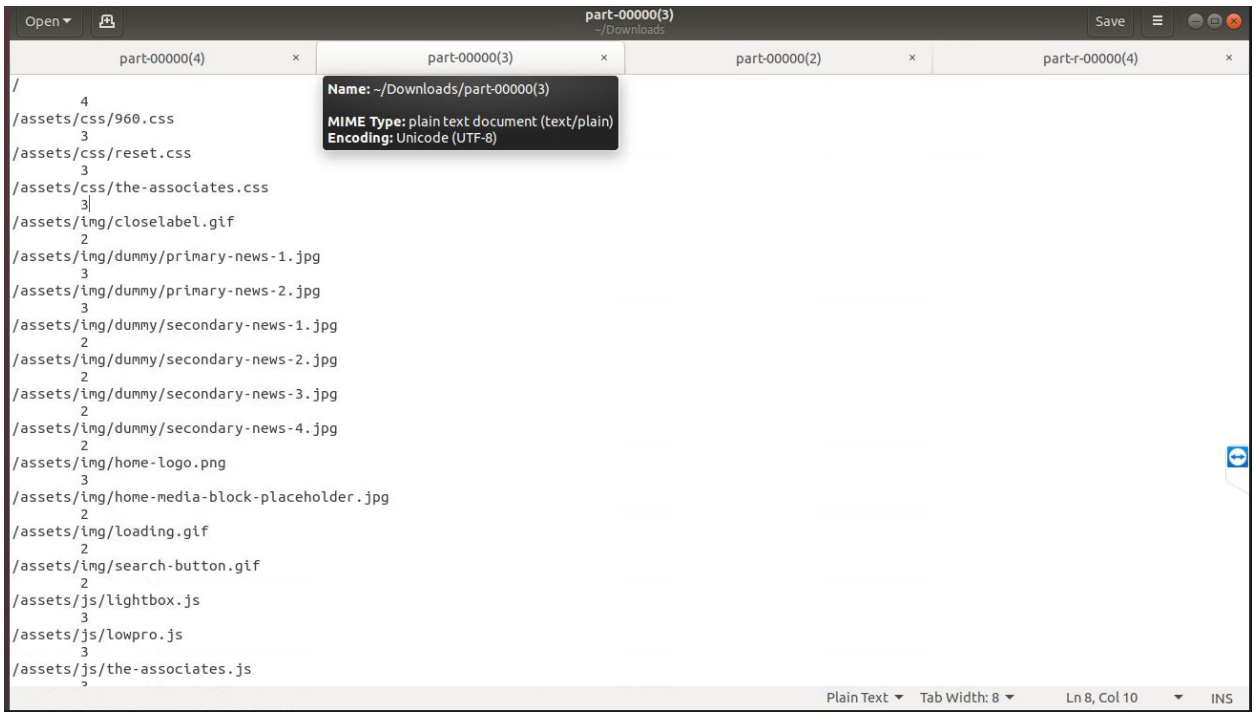

Problem Statement-5: Write a pig script to find max temp from the given dataset.

Description:

160117737006
N.Durga Sai Lakshmi

The Apache Pig MAX function is used to find out the maximum of the numeric values or chararrays in a single-column bag. It requires a preceding GROUP ALL statement for global maximums and a GROUP BY statement for group maximums. However, it ignores the NULL values.

Procedure:

$cd hadoop-3.2.1/

$cd sbin/

$cd start-all.sh

$jps

$cd..

$pig –x local

Code:

grunt>records = LOAD 'sample.txt' AS (year:chararray, temperature:int, quality:int);

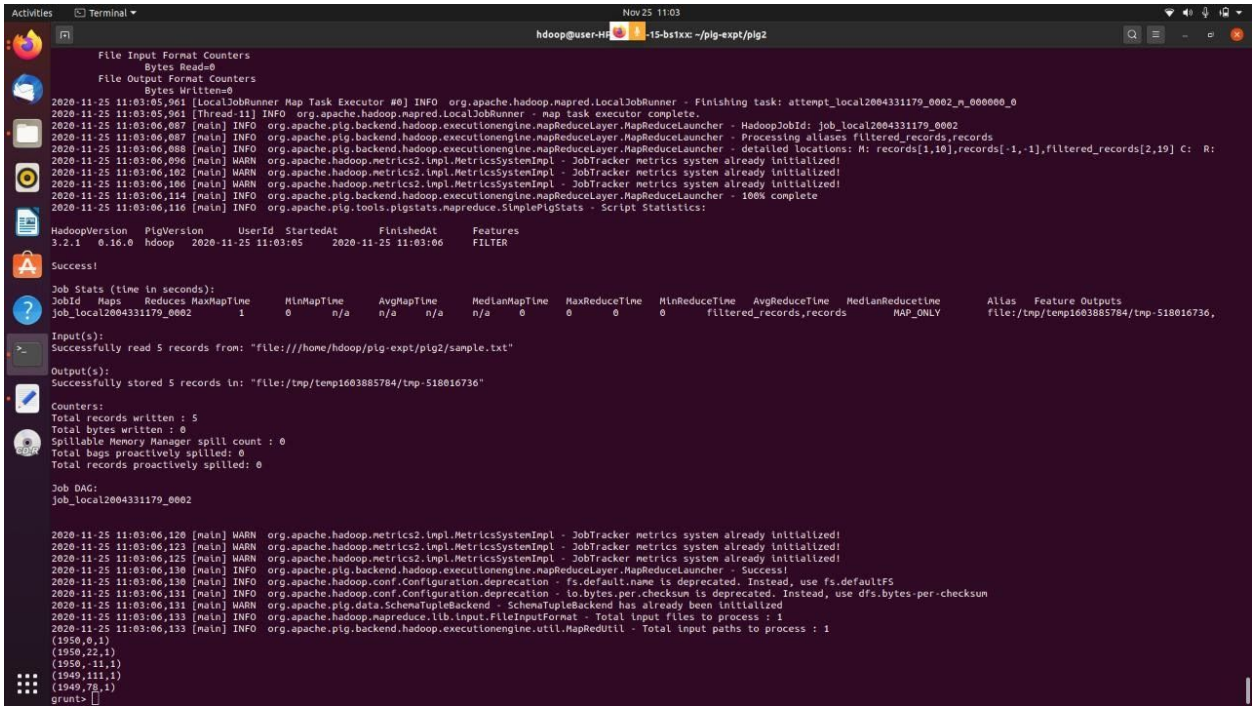

grunt>filtered_records = FILTER records BY temperature != 9999 AND (quality == 0 OR quality == 1 OR quality == 4 OR quality == 5 OR quality == 9);

160117737006
N.Durga Sai Lakshmi

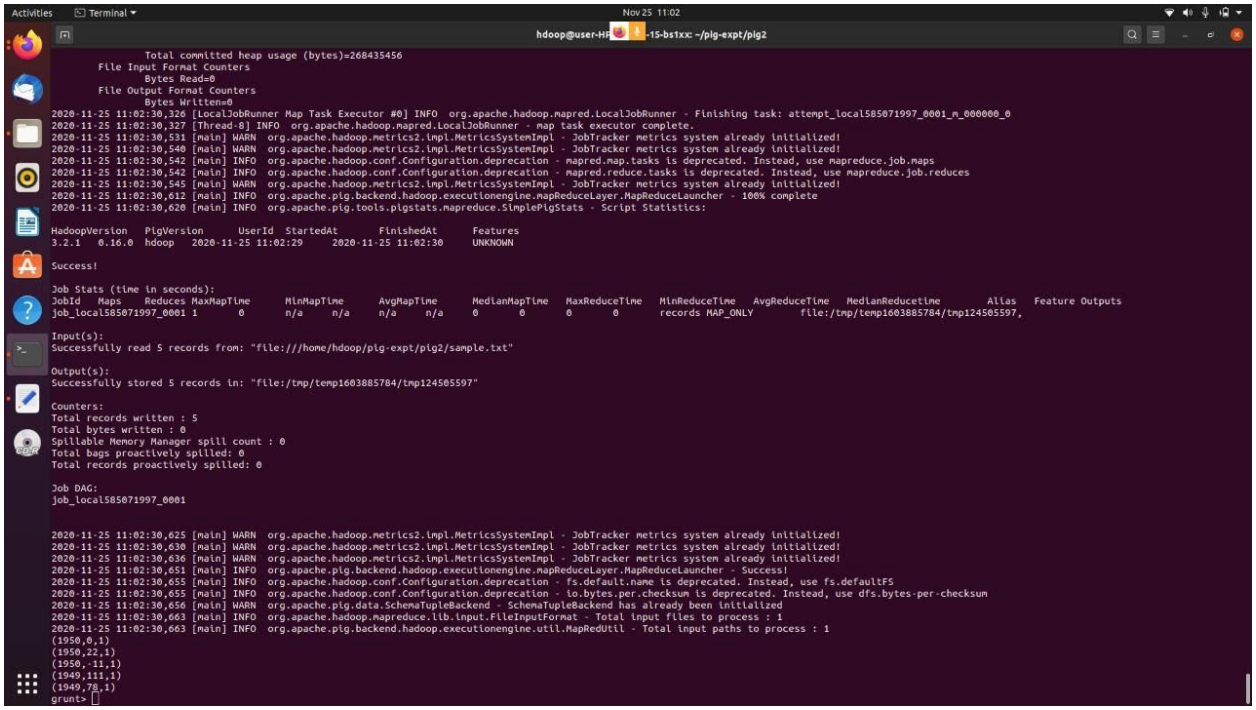

grunt>grouped_records = GROUP filtered_records BY year;

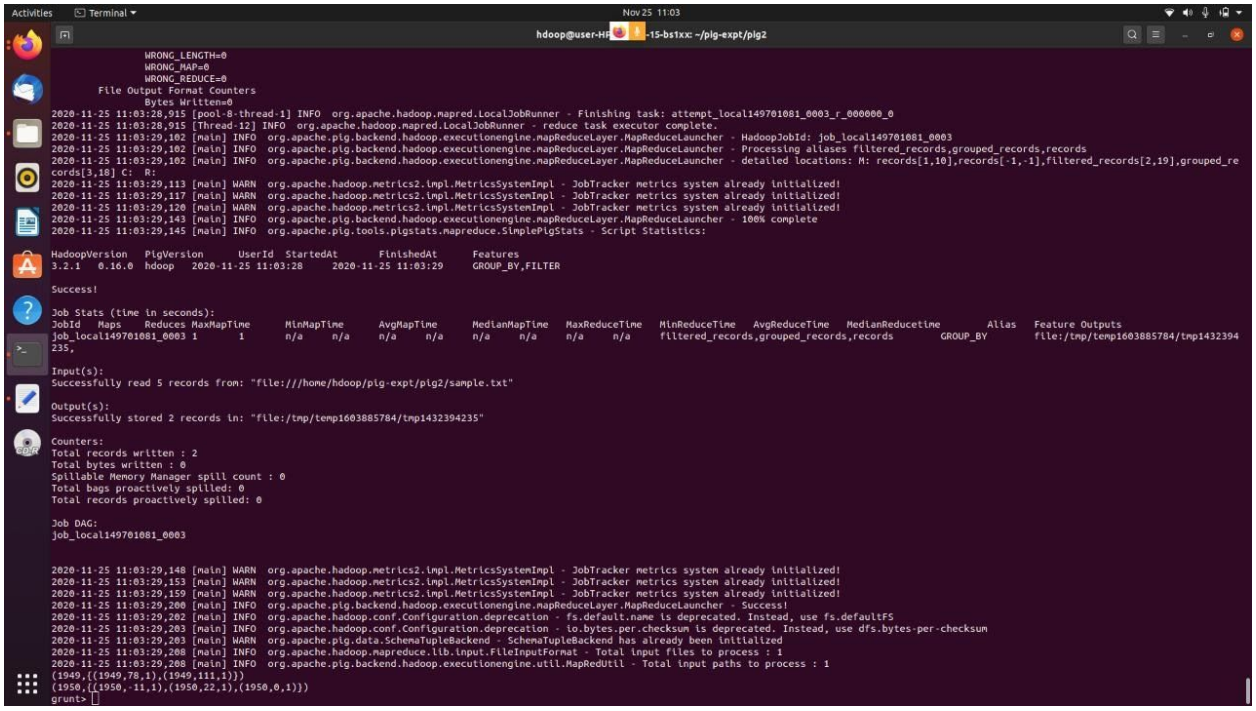

grunt>max_temp = FOREACH grouped_records GENERATE group,
MAX(filtered_records.temperature);

grunt>DUMP max_temp;

160117737006
N.Durga Sai Lakshmi

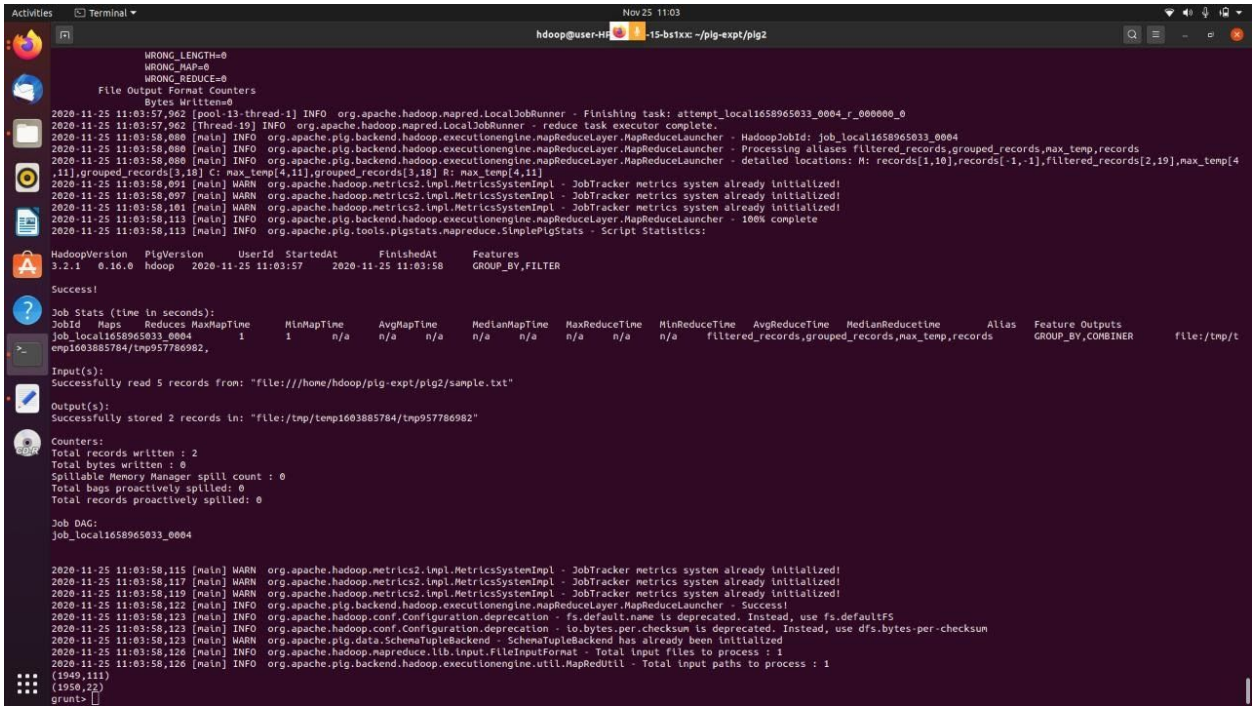

Sample.txt

1950	0	1
1950	22	1
1950	-11	1
1949	111	1
1949	78	1

Problem Statement-6: Write Pig script to implement Word Count Job.

Description: Pig is a high-level programming language useful for analyzing large data sets. Pig was a result of development effort at Yahoo!

In a MapReduce framework, programs need to be translated into a series of Map and Reduce stages. However, this is not a programming model which data analysts are familiar with. So, in order to bridge this gap, an abstraction called Pig was built on top of Hadoop.

Procedure:

$cd hadoop-3.2.1/

$cd sbin/

160117737006
N.Durga Sai Lakshmi

$cd start-all.sh

$jps

$cd..

$pig –x local

Code:

grunt>input1 = load 'sample.txt' as (line);
Output

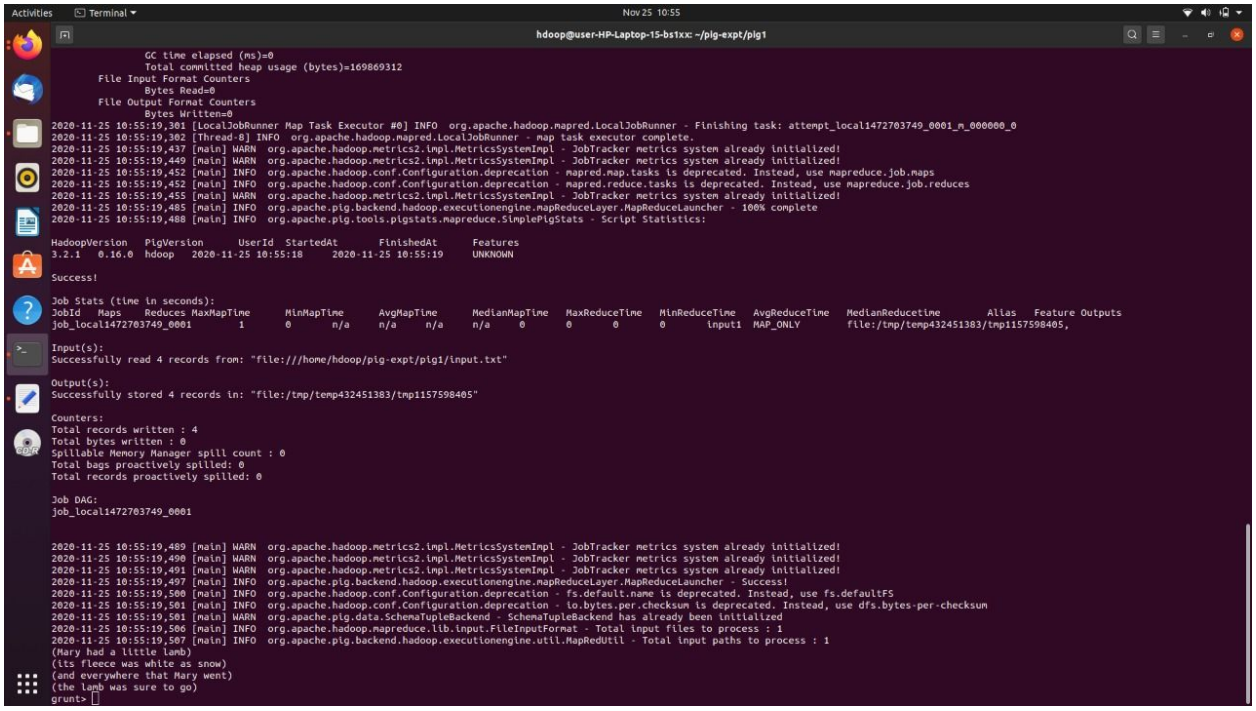

grunt>words = foreach input generate flatten(TOKENIZE(line)) as word;

160117737006
N.Durga Sai Lakshmi

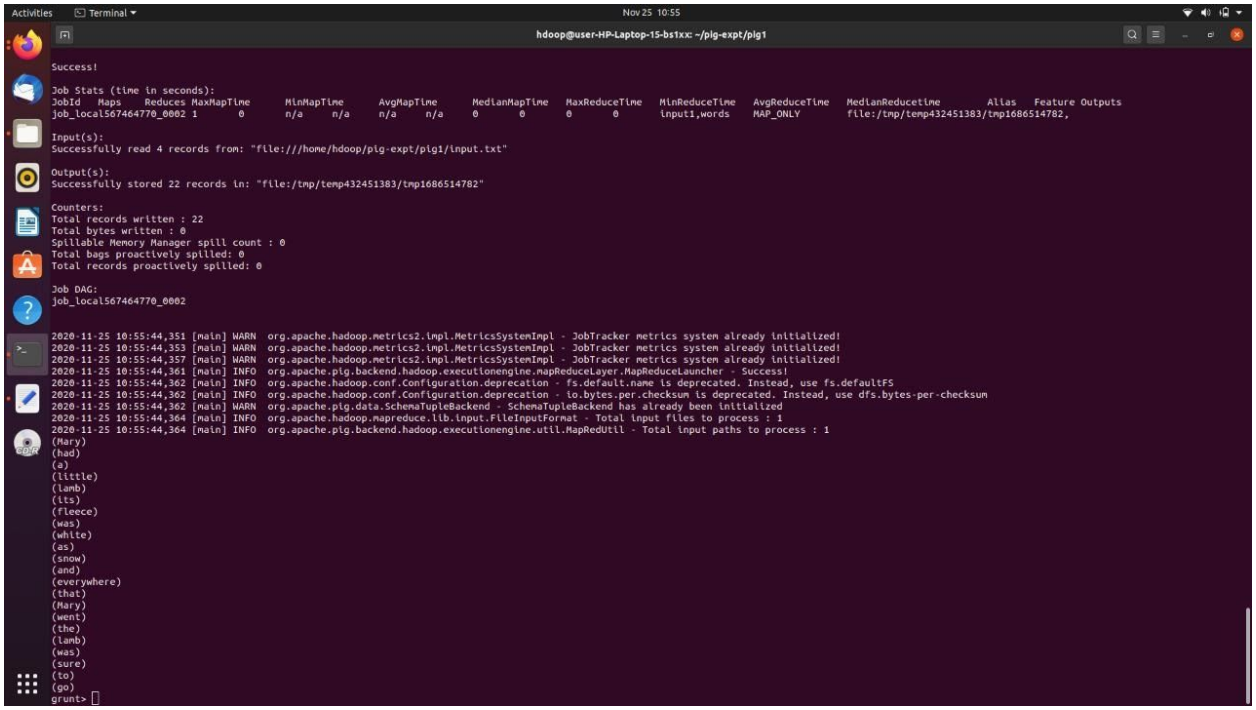

grunt>grpd = group words by word;

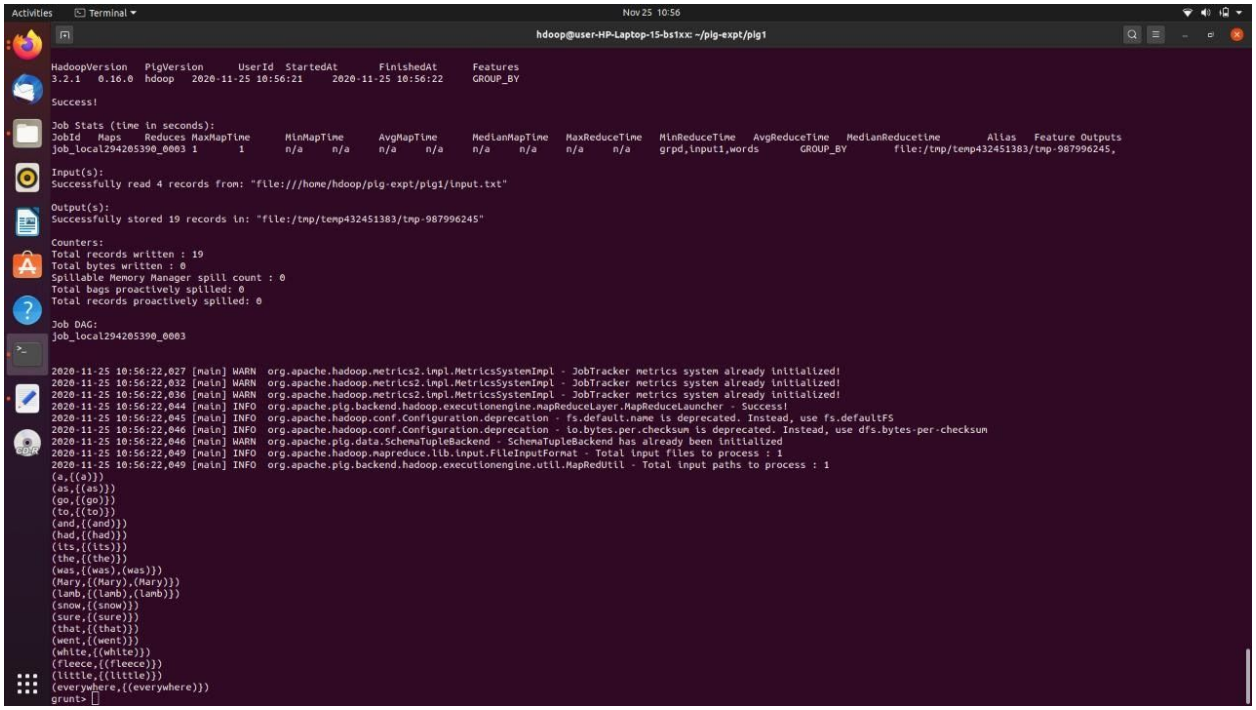

grunt>cntd = foreach grpd generate group, COUNT(words);

grunt>dump cntd;

160117737006
N.Durga Sai Lakshmi

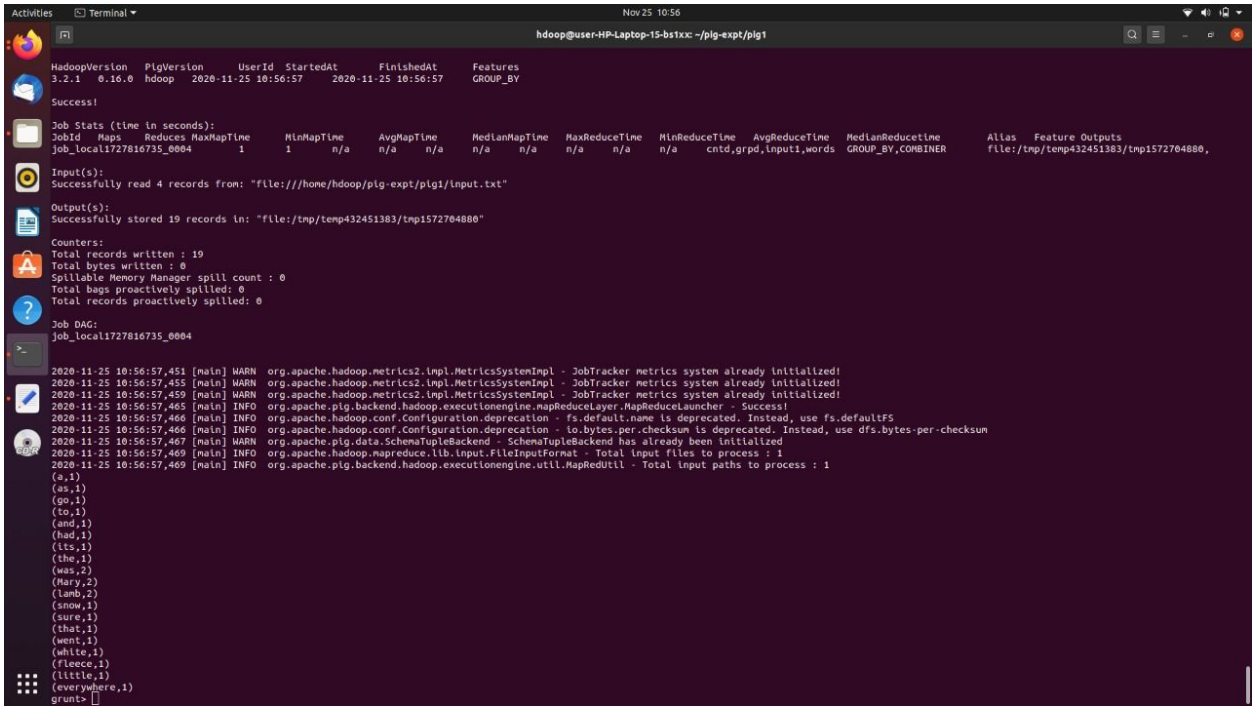

Sample.txt

Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go

Problem Statement-7: Find the Number of Products Sold in Each Country for the given dataset using pig framework.

Description:
Apache Pig enables people to focus more on analyzing bulk data sets and to spend less time writing Map-Reduce programs. Similar to Pigs, who eat anything, the Apache Pig programming language is designed to work upon any kind of data.

Procedure:

$cd hadoop-3.2.1/

$cd sbin/

$cd start-all.sh

160117737006
N.Durga Sai Lakshmi

$jps

$cd..

$pig –x local

Code:

grunt>salesTable = LOAD '/SalesJan2009.csv' USING PigStorage(',') AS (Transaction_date:chararray,Product:chararray,Price:chararray,Payment_Type:chararray,Name:chararray,City:chararray,State:chararray,Country:chararray,Account_Created:chararray,Last_Login:chararray,Latitude:chararray,Longitude:chararray);

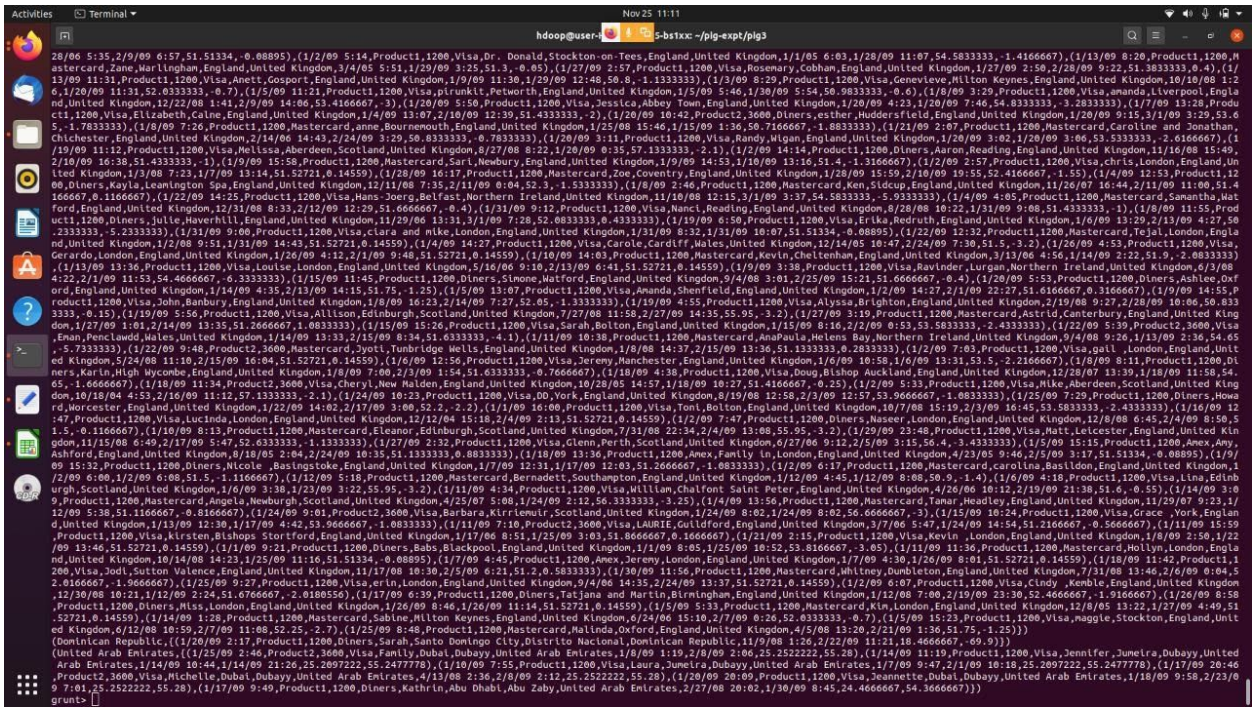

grunt>GroupByCountry = GROUP salesTable BY Country;

160117737006
N.Durga Sai Lakshmi

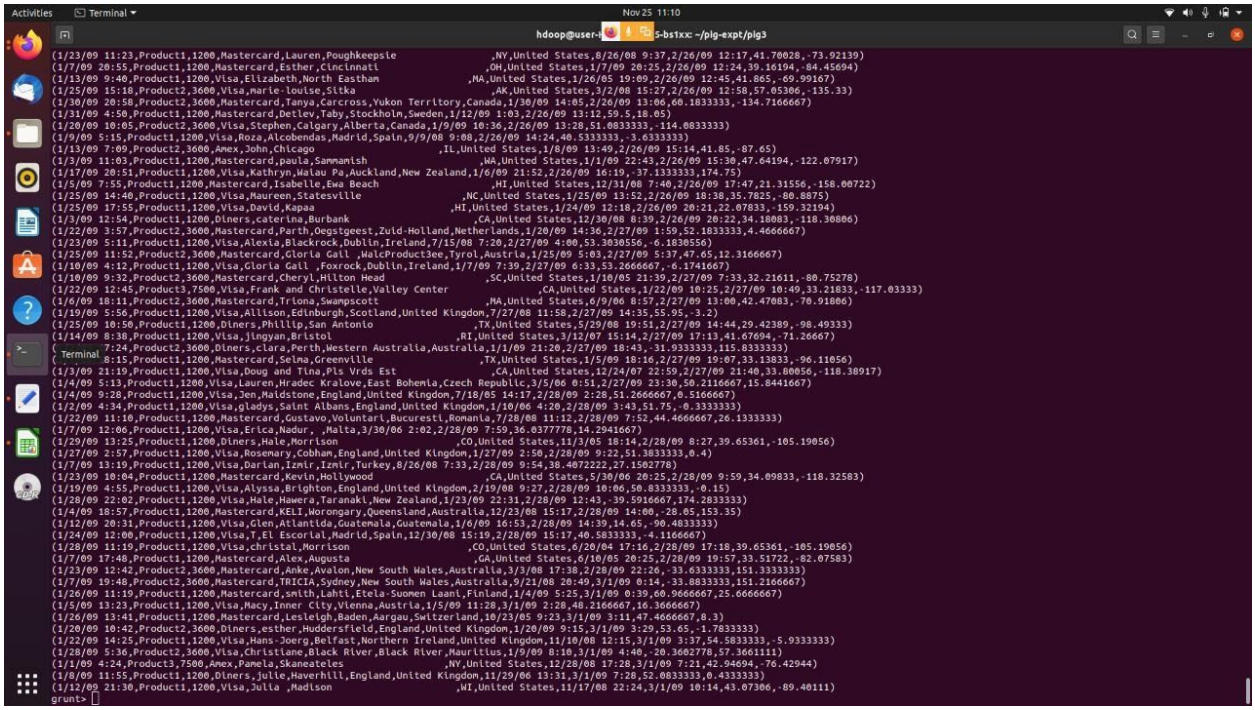


```
grunt>CountByCountry = FOREACH GroupByCountry GENERATE
CONCAT((chararray)$0,CONCAT(':',(chararray)COUNT($1)));
```

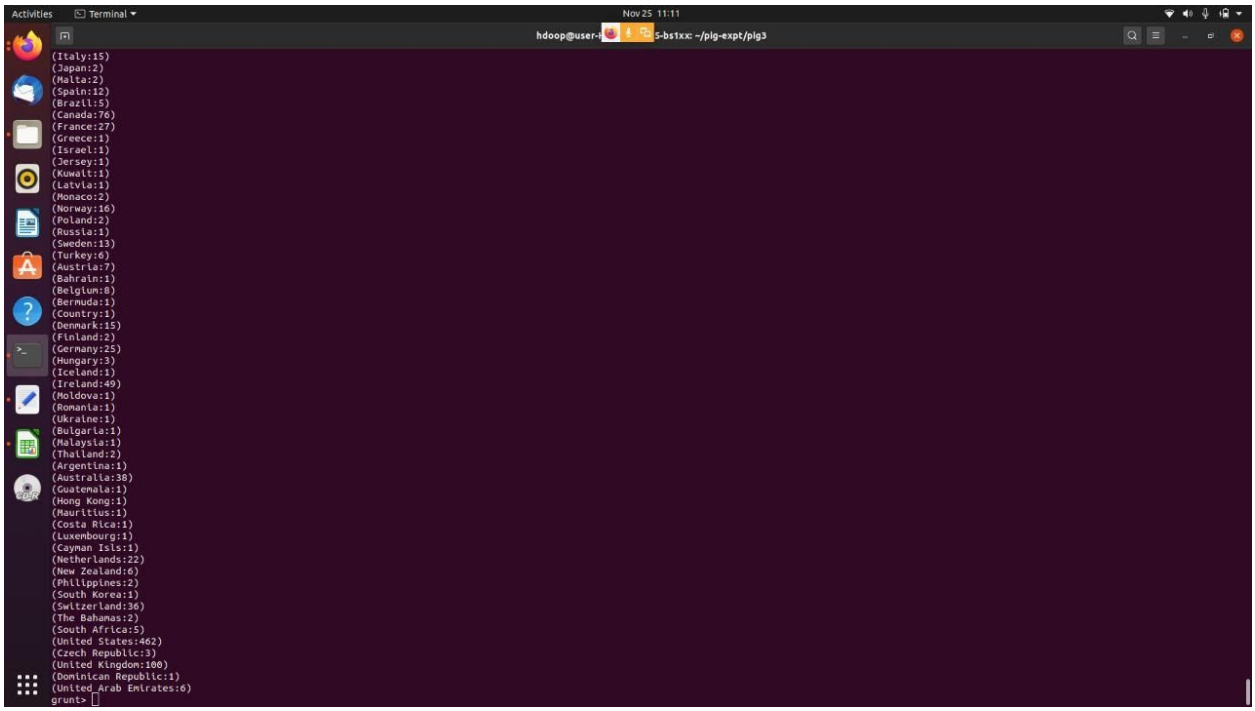


```
grunt>STORE CountByCountry INTO 'pig_output_sales' USING PigStorage('\t');
```

160117737006
N.Durga Sai Lakshmi

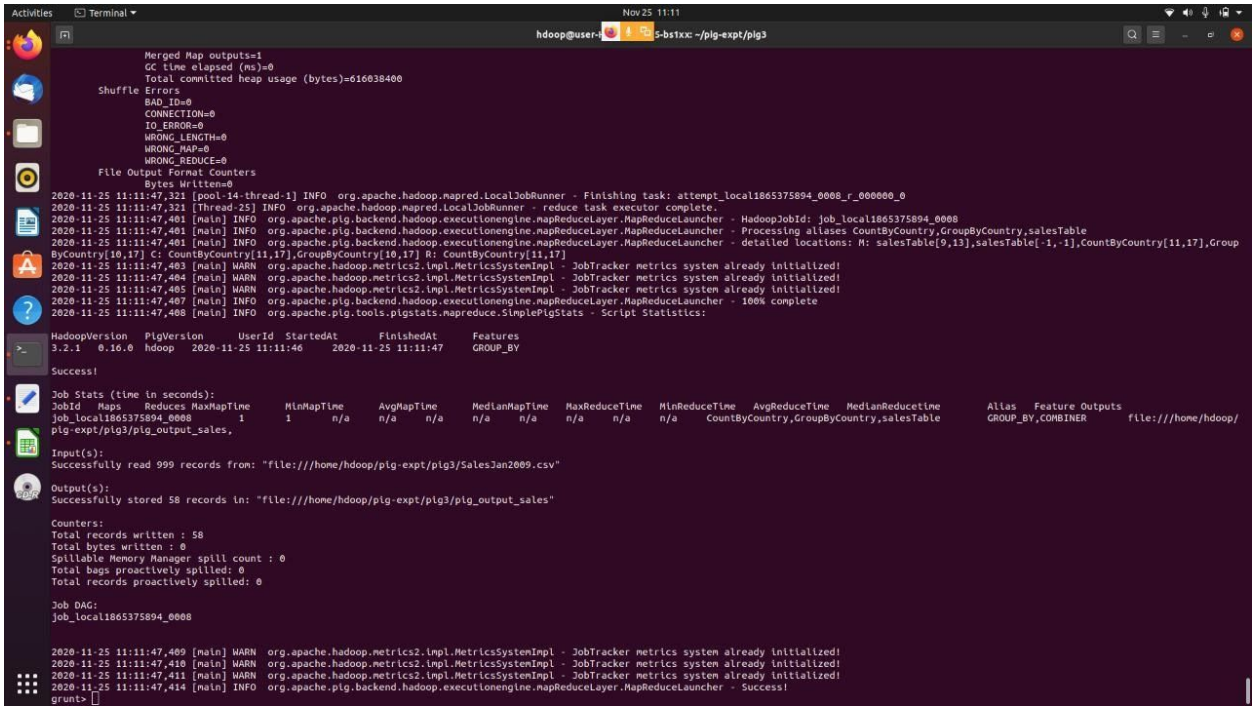

Problem Statement-8: Illustrate the concept of bucketing and partitioning and bucketing in hive.

Description:

Partitioning – Apache Hive organizes tables into partitions for grouping the same type of data together based on a column or partition key. Each table in the hive can have one or more partition keys to identify a particular partition. Using partition we can make it faster to do queries on slices of the data.

The Hive command for Partitioning is:

CREATE TABLE table_name (column1 data_type, column2 data_type) PARTITIONED BY (partition1 data_type, partition2 data_type,….);

Bucketing – In Hive Tables or partitions are subdivided into buckets based on the hash function of a column in the table to give extra structure to the data that may be used for more efficient queries.

The Hive command for Bucketing is:

160117737006

N.Durga Sai Lakshmi

CREATE TABLE table_name PARTITIONED BY (partition1 data_type, partition2 data_type,….) CLUSTERED BY (column_name1, column_name2, …) SORTED BY (column_name [ASC|DESC], …)] INTO num_buckets BUCKETS;

Code:

set hive.exec.dynamic.partition.mode=nonstrict;

Create table students (id int, name string, year int, dept string) row format delimited fields terminated by „,“; Load data inpath „hdfs://localhost:54310/lab10/data“ into table students;

Create table parteddepartment (id int, name string, year int) PARTITIONED by (dept string) row format delimited fields terminated by „,“;

insert overwrite table parteddepartment PARTITION(dept) SELECT id,name,year,dept from students;

Create table samplebucket (id int, name string, year int) clustered by (name) into 2 buckets row format delimited fields terminated by „,“;

From parteddepartment insert overwrite table samplebucket select id,name,year;

```
hive> create table parteddepartment(id int,name string,year int) PARTITIONED BY(dept string);
OK
Time taken: 0.167 seconds
hive> show tables;
OK
parteddepartment
students
Time taken: 0.126 seconds, Fetched: 2 row(s)
hive>
```

```
Loaded : 3/3 partitions.
        Time taken to load dynamic partitions: 0.619 seconds
        Time taken for adding to write entity : 0.001 seconds
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 2.77 sec   HDFS Read: 4272 HDFS Write: 272 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 770 msec
OK
Time taken: 24.837 seconds
```

160117737006
N.Durga Sai Lakshmi

```
hive> show partitions parteddepartment;
OK
dept=cse
dept=eee
dept=it
Time taken: 0.174 seconds, Fetched: 3 row(s)
hive>
```

Problem Statement-9: From the given Aadhaar dataset , write the hive queries for the following.

a. Count the number of Identities generated in each state
b. Count the number of Identities generated by each Enrolment Agency
c. For how many states the Aadhaar data exists?
d. What is the % of Aadhaar being approved per state?

Description:
Hive provides SQL type querying language for the ETL purpose on top of Hadoop file system.

Hive Query language (HiveQL) provides SQL type environment in Hive to work with tables, databases, queries.

We can have a different type of Clauses associated with Hive to perform different type data manipulations and querying. For better connectivity with different nodes outside the environment. HIVE provides JDBC connectivity as well.

Hive queries provides the following features:

- Data modeling such as Creation of databases, tables, etc.
- ETL functionalities such as Extraction, Transformation, and Loading data into tables
- Joins to merge different data tables
- User specific custom scripts for ease of code
- Faster querying tool on top of Hadoop

Code:

create table aad(register string,enrolment_agency string,state string,district string,sub_district string,pincode int,gender string,age int,aad_generated int,enroll_rejected int,res_providing_email int,res_providing_number int)row format delimited fields terminated by ',' stored as textfile;

160117737006
N.Durga Sai Lakshmi

load data local inpath '/home/ak/Desktop/adata.txt' overwrite into table aad;

a. select state,count(*) from aad group by state;

b. select enrolment_agency,count(*) from aad group by enrolment_agency;

c. select count(distinct state) from aad;

d. select state,((sum(aad_generated)-sum(enroll_rejected))/(sum(aad_generated))*100) from aad group by state;

OUTPUTS:

a)

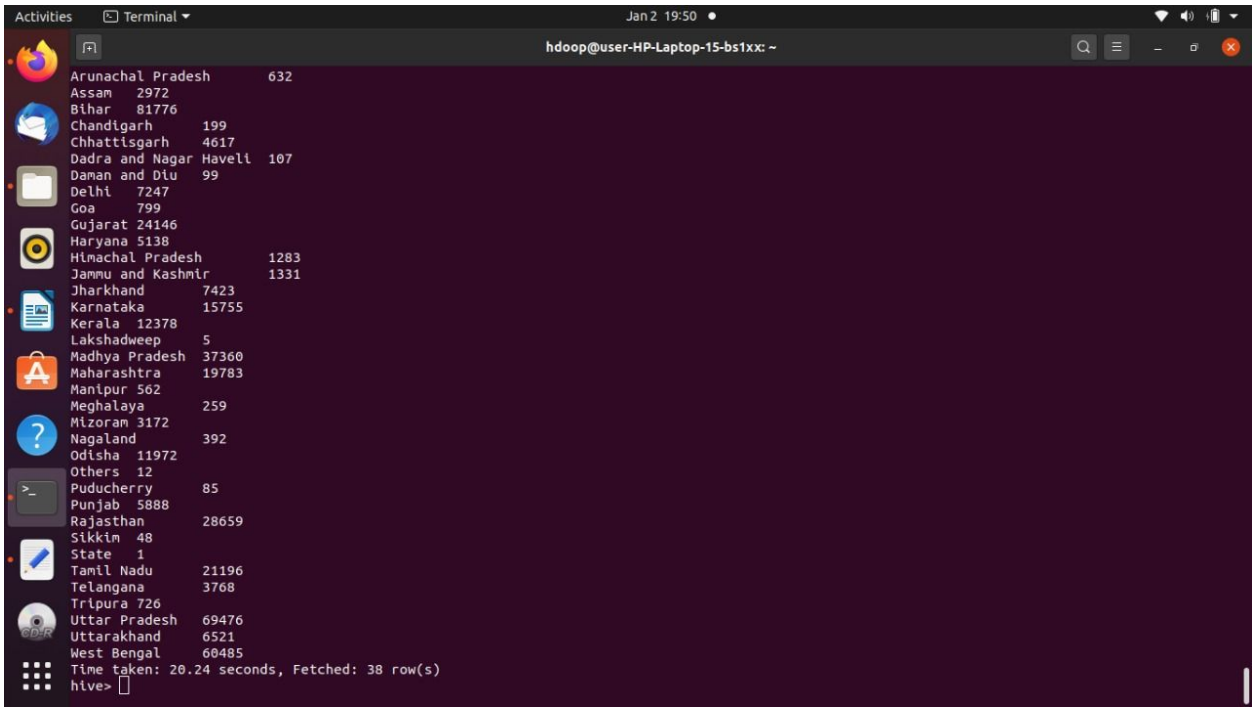

160117737006
N.Durga Sai Lakshmi

b)

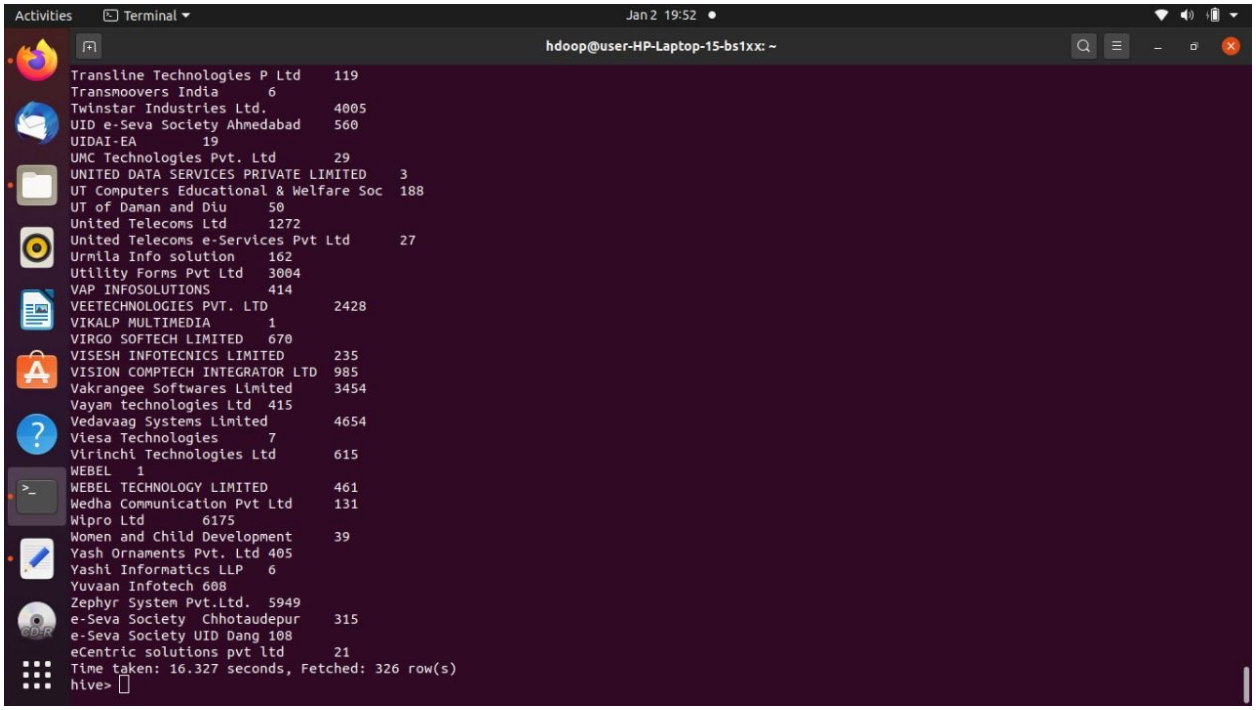

c)

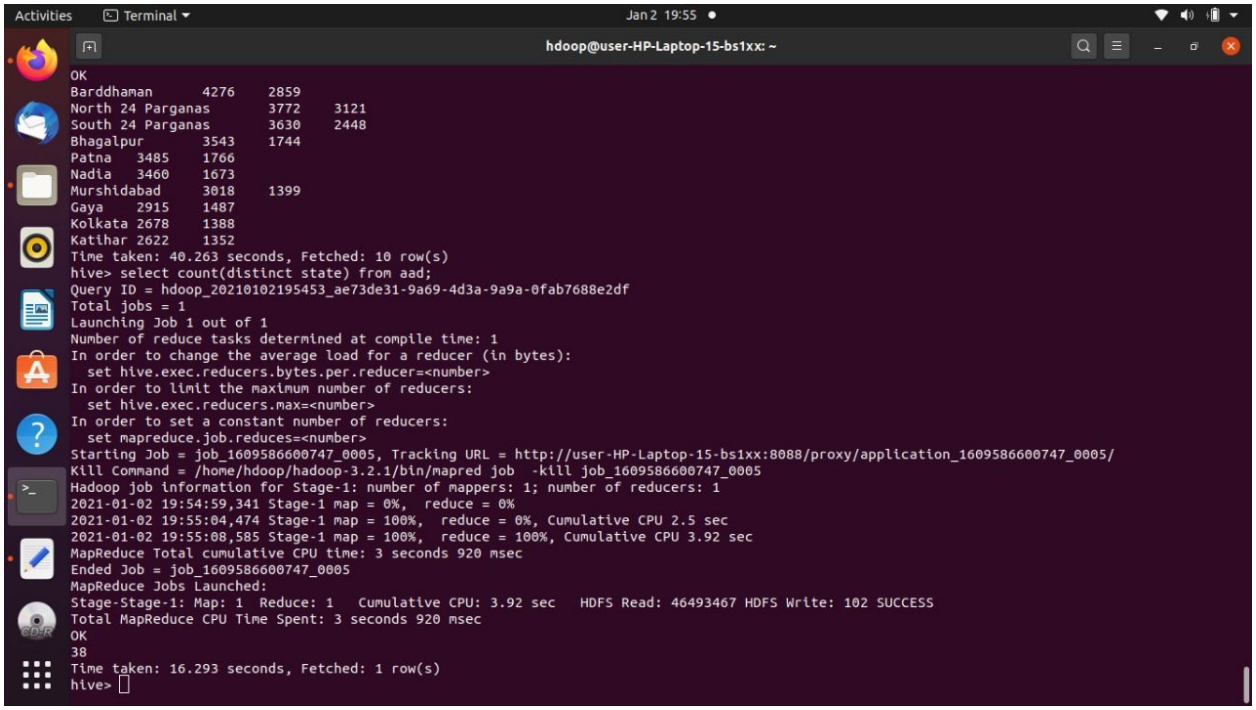

160117737006
N.Durga Sai Lakshmi

d)

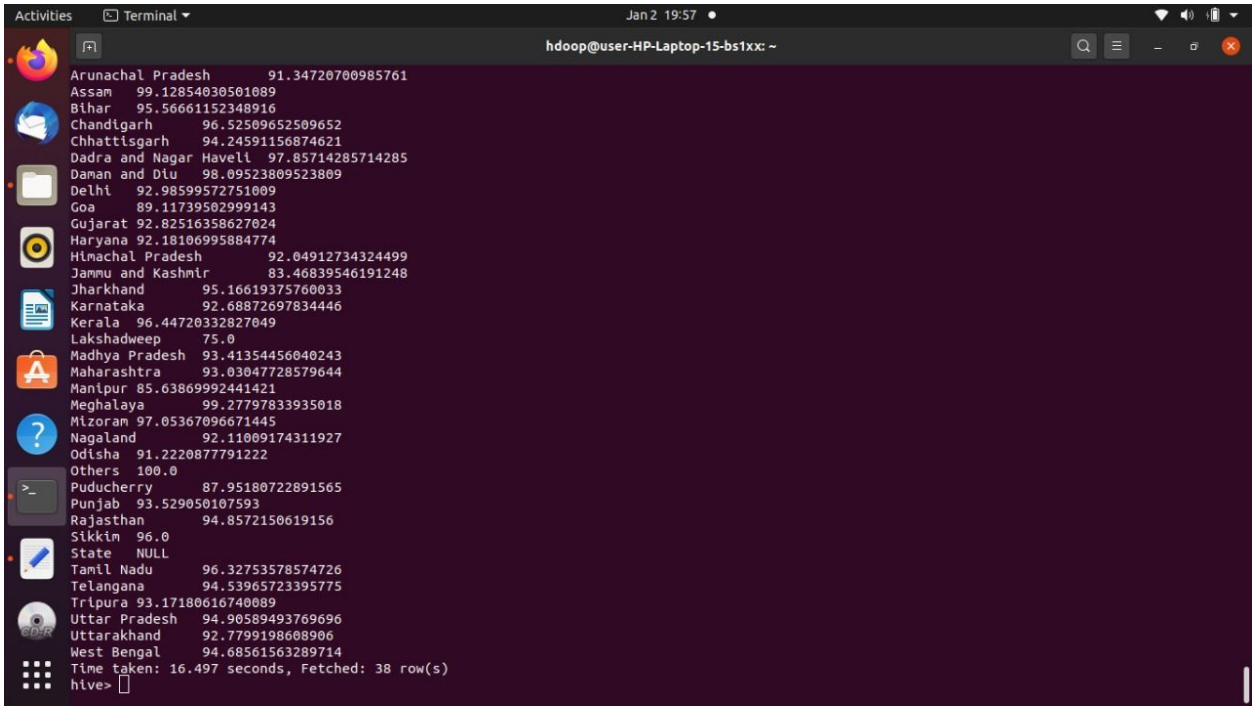

Problem Statement-10: Implement the word count job in Scala.

Description:
Scala is a general-purpose programming language. It supports object oriented, functional and imperative programming approaches. It is a strong static type language. In scala, everything is an object whether it is a function or a number. It does not have a concept of primitive data.

Procedure:
hdfs dfs -mkdir /spark
hdfs dfs -put /home/kgr/sparkdata.txt /spark

Code:
```
scala> val data=sc.textFile("sparkdata.txt");
val splitdata = data.flatMap(line => line.split(""));
val mapdata = splitdata.map(word => (word,1));
val reducedata = mapdata.reduceByKey(_+_);
reducedata.collect;
```

160117737006
N.Durga Sai Lakshmi

For dataset:

```
val text = sc.textFile("mytextfile.txt")
val counts = text.flatMap(line => line.split(" ")).map(word => (word,1)).reduceByKey(_+_)
counts.collect
```

Input file:

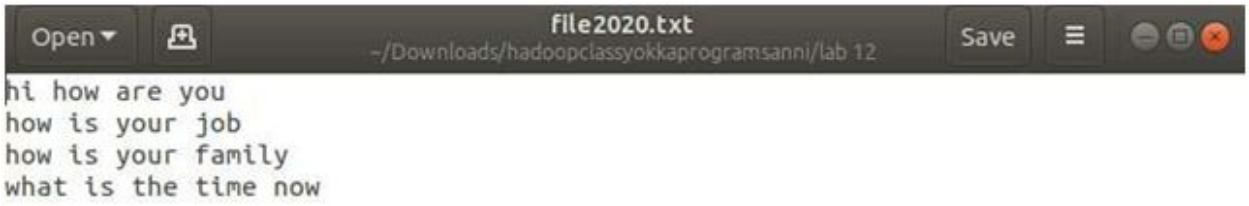

Outputs: